Indexing

Indexing is used to identify the libraries and documents that rules can be applied to. Incremental indexes are run during the day, while a full index is run at night. Indexes check access to every site collection and then process those where:

- SR Enabled is ticked (see Site Collections view Smart Records)
- The iWorkplace Service account has sufficient access
- If incremental - changes have been made since the last index

Auto adding site collections

By default, you must manually select the site collections you want to include in the index and apply rules to. This behaviour can be changed in the settings page to either to automatically include all new sites or use a filter to decide which ones should be added.

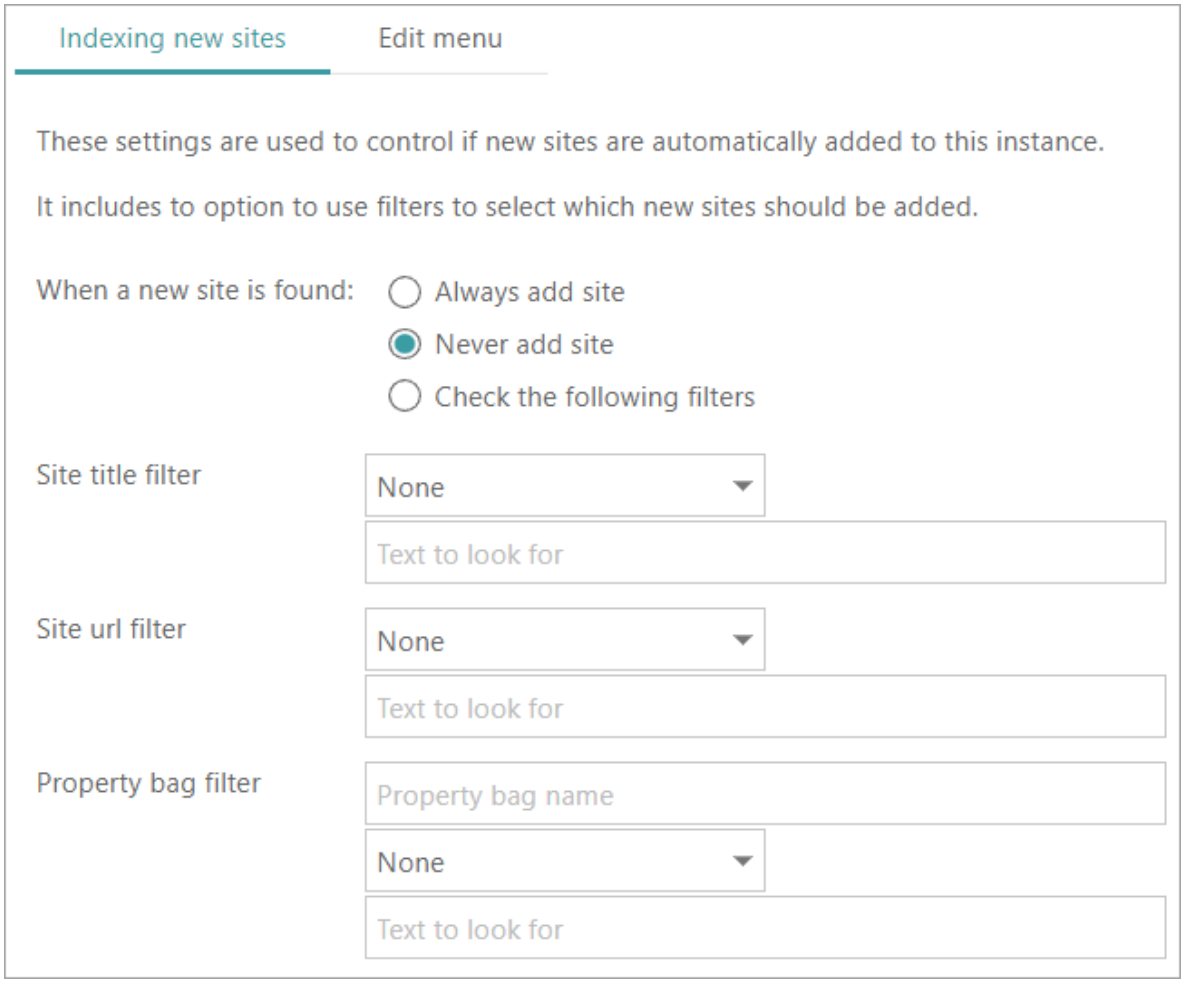

The **Always** and **Never** options match the original functionality, and Smart Records will use existing setting if no changes have been made.

A third option has been added called **Check the following filters** that allows selected adding of new sites. This has three different filters that can be used, either singly or together. When a new site is found, if the filter values match then the site will be added to the instance.

This is not retrospective, so if change is made to these settings then existing sites will not be rechecked.

Site title and url filters

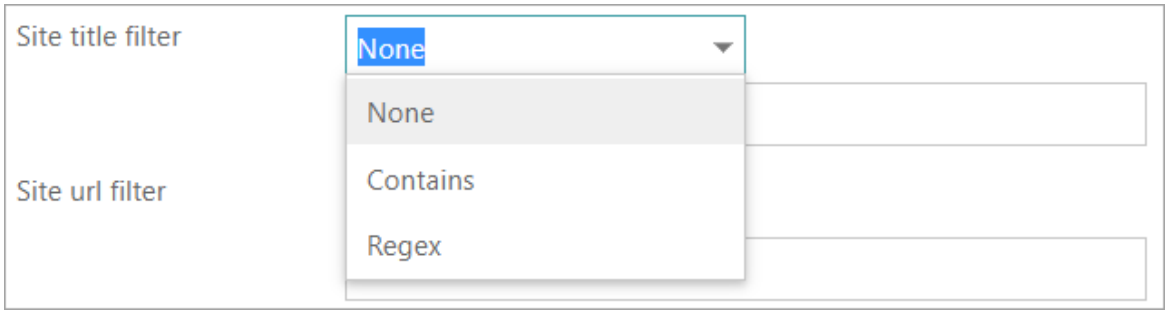

These filters allow you check the Title or Url of the site to see if it matches a particular criterion, for example the title having the word *Project* in it. To use these options first select the type of filter, then enter the value. Filter types available are:

- **None** – this must always be selected if you do not want to use this filter type. The reason for this is that selecting Contains or Regex and then leaving the value blank can still end up in the items being filtered.
- **Contains** – does the title or url contain a certain value. The check is not case sensitive so, for example, both *PROJECT* and *project* would match. This is a simple contains match and you cannot use additional wild card characters in the value, for that sort of matching use the Regex option.
- **Regex** – this is for complex pattern matching using regex formulas.

Property bag filter

This option lets you look for a particular property bag in the site. If it exists, then the value will be checked to see if it matches. This has the following filter types:

- **None** – this must always be selected if you do not want to use this filter type.
- **Equals** – does the property bag value match the entered value. This check is not case sensitive.
- **Contains** – does the property bag value include the entered value. The check is not case sensitive.
- **Starts with** – does the property bag value start with the entered value. The check is not case sensitive.

There is no regex option with the property bag check.

When the index runs…

The first step the index does is to check the column information to see if there have been any changes that need to be reflected in the new index e.g. a new managed column has been defined.

The indexer goes through each site collection and looks for the ones that have the SR Enable ticked. If this is ticked it checks if the iWorkplace service account has sufficient access to the site collection.

If both are true then it iterates through each subsite, library and document and update the index accordingly. The Site Collections view in Smart Records lets you configure any site collection so the iWorkplace Service account has access, you can also tick the SR Enable box for any site.

Incremental indexes

Incremental indexes only process sites, libraries and documents that have been added, modified, or deleted since the last full or incremental index. An incremental index is run each night, the exact time will vary as the process uses random delays etc to prevent throttling.

Full indexes

The indexing process is designed to always run an incremental index unless:

1. It detects a new column to index has been added in Smart Records
2. Or new functionality has been added to Smart Records that requires a full index
3. Or the index failed part way through due to system outage

If any of these occur, then a full rather than incremental index will be run. Full indexes can take a long time to run as every entry is checked and updated as required. Run times vary depending on the amount of data, number of

changes since the last index, etc to its hard to estimate but typically 4-6 hours for small to medium setups and up to 48 hours for large systems.

Licence limit check

Once complete the number of records in the index is compared with the licence limit. If the number of records is over 90% of the licence limit, then the Records and Licence Limit widgets will turn orange. If the licence limit has been exceeded, the records data will be truncated to the licence limit – so data on some records will not be available in Smart Records. If this happens both the Records and Licence Limit widgets on the home page will turn red.

Indexer summary

The SR365 homepage shows a summary of the indexer start, finish and the number of items, libraries and sites that have been indexed.

Index Summary

Full Index Start	2015-12-10 22:00:00
Full Index Finish	2015-12-10 23:04:42
Incremental Start	2015-12-10 13:00:00
Incremental Finish	2015-12-10 13:45:43
Record Count	1,775
Library Count	224
Site Count	91
Site Collection Count	40
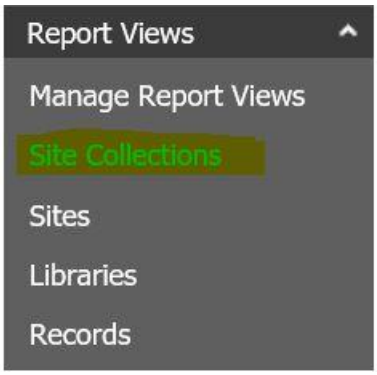
You can also see a breakdown of when the indexer last ran, by site collection by selecting the "Site Collections" link under the Report Views: